

Lab5 Report

Jiangbei Li

- Design Decision
 - LockManager
 - In the class, I create 4 maps to save the locks
 - Map1&2: PageId -> TransactionId (Read&Write)
 - Map3&4: TransactionId -> PageId (Read&Write)
 - I also create two maps to save the query that waits lock
 - Map5&6: PageId -> TransactionId (Read&Write)
 - I create another map to save a mutex for each page to keep synchronized:
Map7: PageId -> mutex
 - I use concurrentHashMap instead of HashMap and HashSet.
 - Function: grantLock
 - Use synchronized for the mutex of the page.
 - If there exist ww, wr, rw, then add the query to wait map; else add lock to the tid and pid.
 - Function: releaseTidLock
 - Release all the locks that the TransactionId holds and the waiting query of the TransactionId
 - Function: releaseLock
 - Release the lock for the tid and pid.
 - Function: holdsWriteLock
 - Return whether the TransactionId hold write lock for the page.
 - Function: holdsLock
 - Return whether the TransactionId hold lock for the page.
 - Function: deadlockOccur
 - Detect the deadlock for the buffer
 - BufferPool

- Add class member lockManager to keep lock.
- Function: getPage
 - Call grantLock to request a lock for the query
 - If not grantLock, then call deadLockOccur to detect deadLock and keep calling grantLock until granted
- Function: releasePage and holdsLock
 - Call the corresponding function in lockManager.
- Function: transactionComplete
 - If the commit is True: For all pages that the tid has write lock on it, I flush these pages and mark them not dirty.
 - If the commit is not True: I call getBeforeImage and fetch the old page from Disk to buffer.
- Deadlocks
 - I create directed dependency graph for each transaction and page
 - For one transaction, if it has a lock, and another transaction is waiting for the lock (ww, wr, rw), then there is an edge from the first transaction to the second.
 - From the input transaction, I use BFS(Breadth first search) to traverse the graph: if there is a cycle for the start vector then there will be a deadLock.
- Difficulty and time
 - I spend four days on this lab
 - I use hashSet firstly but it is not concurrent, then I use concurrentHashMap and use the keyset to replace hashSet.
 - I forgot to write writePage in HeapFile. I comment on a necessary code in insertTuple in HeapPage. Then the inserted tuple will have no recordId
 - When running TransactionTest, I catch the TransactionAbortedException ahead of the TransactionTestClass. Then the query will not rerun.
 - The BTreeTest is very hard and I cannot find a good way to debug.
- Weakness
 - The running time is very long for BTreeTest and may fail with a very small possibility. I am trying to make it faster.